



ILLINOIS TECH

Illinois Institute of Technology

CSP 554 – Big Data Technologies

Project Report

Electricity Consumption Forecasting using Big Data and Deep Learning

Date: Dec 10, 2025

Team Members

Name	Hawk ID
Chaitanya Datta Maddukuri	A20568393
Tanushree Sharma	A20577521

Submitted to:

Prof. Joseph Rosen
Department of Computer Science
Illinois Institute of Technology

Table of Contents

- 1. Abstract**
- 2. Introduction**
- 3. Objective**
- 4. Scope**
- 5. Dataset Overview**
- 6. Significance**
- 7. Literature Review**
- 8. Data Preprocessing**
- 9. Tools Used**
- 10. Exploratory Data Analysis (EDA)**
- 11. Predictive Modeling**
- 12. Big Data Aspect**
- 13. Results**
- 14. Challenges Faced**
- 15. Conclusion and Future Work**
- 16. Acknowledgment**

References

1 Abstract

Short-term electricity load forecasting is a critical task for power system operators, balancing authorities, and market participants. With the widespread deployment of smart meters and advanced metering infrastructure, utilities now collect massive volumes of high-frequency consumption data, often enriched with weather and calendar information. Traditional statistical models struggle to handle this combination of **volume, variety, and temporal complexity**, motivating the use of both **big data technologies** and **deep learning**.

This project presents an end-to-end forecasting pipeline that integrates **Amazon S3, Apache Spark on Amazon EMR, and deep learning models (LSTM and GRU) executed on EC2** for **next-hour electricity load prediction**. Hourly load data from PJM-like regions is stored in S3 as partitioned Parquet files and processed using PySpark for cleaning, feature engineering, and exploratory data analysis (EDA). Weather features (temperature, humidity, wind speed), temporal indicators (hour, day-of-week, month), and engineered lag/rolling features are used as inputs to deep sequence models.

Due to capacity and quota limitations for SageMaker Studio notebook instances, model training was performed directly on the EMR master node / EC2 rather than through SageMaker. Despite these constraints, optimized LSTM and GRU networks achieved strong predictive performance, with the LSTM model reaching approximately **MAE $\approx$ 550 MW** and **RMSE $\approx$ 691 MW** on the held-out test set. The project demonstrates that a combination of **cloud-native big data tools** and **deep learning architectures** offers an effective and scalable approach for real-world electricity load forecasting.

2 Introduction

Electricity demand varies over time due to a wide range of factors, including human activity patterns, weather conditions, holidays, and economic conditions. System operators must continuously balance generation and consumption; failure to do so can lead to frequency deviations, equipment damage, or even blackouts. **Short-term load forecasting (STLF)** typically in the range of one hour to several days ahead—is therefore a fundamental function in power systems.

Historically, STLF has been addressed with statistical models such as ARIMA, exponential smoothing, and regression-based methods. While these techniques perform reasonably well for smooth, relatively linear time series [2], they often fail to capture the nonlinear effects and complex interactions inherent in modern electricity consumption data. As electric grids grow more dynamic, with demand response, electric vehicles, and distributed energy resources, the underlying data becomes richer but also more challenging.

At the same time, the scale of available data has exploded. Smart meters report consumption at hourly or sub-hourly resolution for millions of customers, generating **big data** both in volume and velocity. Traditional single-machine data processing and modeling pipelines are not sufficient to handle such datasets efficiently; instead, **distributed frameworks like Apache Spark** are needed for ETL (Extract–Transform–Load) and feature engineering.

This project sits at the intersection of **big data processing** and **deep learning**. It aims to design and implement a scalable forecasting pipeline that:

- Reads and processes large amounts of historical load and weather data stored in Amazon S3.
- Uses Apache Spark on EMR for distributed data cleaning, feature engineering, and EDA.
- Trains **Long Short-Term Memory (LSTM)** and **Gated Recurrent Unit (GRU)** models, which are well suited to capturing temporal dependencies in time series.
- Evaluates model performance using standard forecasting metrics.

The outcome is not just a single model, but a **reusable architecture** that can be extended or adapted for other time-series forecasting tasks in a big-data environment.

3 Objective

The objectives of this project can be grouped into **technical**, **analytical**, and **practical** goals.

Technical Objectives

1. **Build a cloud-based data pipeline** that uses Amazon S3 as a data lake and Apache Spark on EMR for distributed processing of multi-year hourly load data.
2. **Implement feature engineering at scale**, including lag features, rolling window statistics, and weather interactions, using PySpark.
3. **Implement deep learning models (LSTM and GRU)** for short-term load forecasting using TensorFlow/Keras on EC2 within the EMR environment.

Analytical Objectives

1. Perform **exploratory data analysis (EDA)** to understand temporal patterns (hourly, daily, weekly, seasonal) in electricity load.
2. Analyze the **relationship between load and weather** variables such as temperature, humidity, and wind speed.
3. Evaluate model performance using **Mean Absolute Error (MAE)** and **Root Mean Squared Error (RMSE)** and compare LSTM vs GRU.

Practical Objectives

1. Demonstrate that it is feasible to handle a **large-scale, real-world dataset** with big-data tools within course constraints.
2. Adapt the architecture to **resource limitations**, by using EMR + EC2 for training when SageMaker notebook capacity is insufficient.
3. Produce a **well-documented and reproducible workflow** that can serve as a template for future big-data forecasting projects.

4 Scope

The scope of this project is intentionally focused, so that the core pipeline is completed and evaluated with reasonable depth.

Included in Scope

- **Data ingestion and storage** of preprocessed load features in Amazon S3.
- **Distributed preprocessing and feature engineering** using Apache Spark on EMR.
- **EDA** focusing on:
 - Temporal patterns (hour-of-day, day-of-week, month).
 - Region-wise differences.
 - Correlations between load and weather.
- **Supervised learning models** for **one-step-ahead** (1-hour) load forecasting.
- **Deep learning models**:
 - LSTM-based sequence model.
 - GRU-based sequence model.
- **Model evaluation** and comparison using MAE and RMSE.
- **Documentation** of architecture, methods, and results.

Out of Scope

- Multi-step-ahead forecasting (e.g., predicting 24 hours ahead in a single model).
- Real-time production deployment (e.g., APIs, streaming inference).
- Fully managed SageMaker training and endpoints (not feasible given instance capacity limits).
- Hardware acceleration with GPUs (training performed on CPU-based EC2 in EMR for this project).

This scoped approach allows the project to **go deep on one-hour-ahead prediction**, while leaving natural extensions for future work.

5 Dataset Overview

The electricity load dataset [7-8] used in this project is derived from hourly historical records of regional power systems similar to those operated by **PJM Interconnection**. The data includes:

- **Timestamp**: indicating the date and hour of the measurement.
- **Region**: such as AEP, COMED, DAYTON, PJM, etc.
- **Load in MW**: the total power demand in megawatts for that region and hour.

These kinds of datasets are widely used in academic and industrial forecasts due to their real-world relevance and granular time resolution.

Weather Data

To capture exogenous drivers of demand, the project incorporates **weather variables** aligned to the same timestamps:

- **Temperature, Humidity, Wind speed**

These features are usually obtained from datasets similar to NOAA's Global Historical Climatology Network (GHCN), and then merged with load data based on time and approximate location.

Storage and Partitioning

All processed and feature-engineered data is stored in **Amazon S3** under the bucket path:

- `s3://csp554-electricityforecast/features/features/`

Within that bucket, data is partitioned by **region**, **year**, and **month**, for example:

- `region=AEP/year=2012/month=01/part-*.parquet`

Partitioning provides two key benefits:

1. **Performance**: Spark can perform predicate pushdown, scanning only relevant partitions for a given query (e.g., a specific region and year).
2. **Organization**: The data is logically grouped, making it easier to run experiments for specific regions or time windows.

Data Volume and Characteristics

Across all regions and years, the dataset contains:

- **Millions of hourly records**, covering multiple years.
- A diverse range of operating conditions: peak vs off-peak, seasonal extremes, holidays.
- Significant missingness in some weather variables (e.g., hundreds of thousands of nulls across temperature, humidity, and wind speed).

These characteristics make the data sufficiently large and complex to benefit from big data processing.

6 Significance

The significance of this work can be examined from **energy system**, **technical**, and **educational** perspectives.

Energy System Perspective

For utilities and system operators:

- Better short-term load forecasts directly reduce **operating costs**, because generation schedules can be planned more efficiently.
- Accurate forecasts reduce the need for expensive **spinning reserves** and lower the risk of **load shedding**.
- Improved forecasting supports **integration of renewables**, whose intermittency adds uncertainty to net load.

Even a small reduction in forecasting error can translate into significant financial savings and improved reliability for large power systems.

Technical Perspective

From a technical point of view, this project is significant because it demonstrates how to:

- Combine **distributed data processing (Spark)** with **deep learning models** in a coherent architecture.
- Handle **large datasets** in a way that is scalable and cloud-native.
- Use advanced sequence models (LSTM/GRU) rather than treating time-series as static tabular data.

This aligns with modern industry practice where **data engineering and machine learning are tightly integrated**.

Educational Perspective

For the CSP 554 course, this project:

- Provides a concrete example of “**Big Data Technologies**” applied to a meaningful, real-world problem.
- Helps understand how concepts like **EMR clusters**, **S3 data lakes**, and **Spark jobs** come together in a full project.
- Reinforces the importance of end-to-end thinking: not just building a model, but managing data, preprocessing, training, evaluation, and interpretation.

7 Literature Review

The literature on load forecasting spans several generations of methods [5-6]. Here we highlight the most relevant streams and how they relate to our approach.

Classical Time-Series Methods

Traditional STLF methods include:

- **ARIMA / SARIMA models** that capture autoregressive and moving-average behavior with seasonal components.
- **Exponential smoothing** for trend and seasonality modeling.

These methods are well studied but assume relatively stationary time-series and linear relationships, which limits their effectiveness when nonlinear effects (e.g., weather interactions) are strong.

Machine Learning and Deep Learning Models

Apache Spark MLlib enables distributed machine learning and offers scalable algorithms for regression, decision trees, and gradient boosting suitable for large datasets [10].

Machine learning brought more flexibility:

- **Random Forests** and **Gradient Boosted Trees** can approximate nonlinear functions and interaction effects.
- **Support Vector Regression (SVR)** can handle some nonlinearity with kernel tricks.

These models, however, generally treat lagged variables as static features and do not have an explicit notion of sequence or memory. They can perform quite well but may miss long-range temporal dependencies.

TensorFlow provides deep learning infrastructure capable of modeling complex temporal patterns using architectures such as LSTM and GRU [11]. AWS SageMaker offers a managed cloud service for scalable model training, tuning, deployment, and monitoring, and includes advanced forecasting algorithms such as DeepAR. Together, these technologies enable both batch and real-time forecasting workflows in modern power systems.

Deep learning introduced **Recurrent Neural Networks (RNNs)**, especially:

- **Long Short-Term Memory (LSTM)** networks, which maintain explicit memory cells controlled by input, output, and forget gates. They are particularly effective for long sequences and avoid vanishing gradient issues.
- **Gated Recurrent Units (GRU)**, a simplified variant that uses fewer gates but can achieve similar performance.

Research shows that Long Short-Term Memory (LSTM) and Gated Recurrent Unit (GRU) architectures outperform many classical forecasting models on multivariate time-series datasets due to their ability to capture long-term temporal dependencies. Hybrid approaches such as CNN-LSTM networks and

sequence-to-sequence architectures have also been studied to enhance short-term energy consumption forecasting

Several works in the literature demonstrate that LSTM-based models outperform ARIMA, traditional ML, and basic RNN [3] for short-term load forecasting when sufficient data and features are available. They can directly ingest sequences of lagged load, weather, and calendar information.

Big Data and Cloud-Based Forecasting

As electricity datasets continue to grow, researchers and industry practitioners increasingly use:

- **Hadoop and Spark** [9] for distributed data processing.
- **Cloud platforms (AWS, Azure, GCP)** [12] to handle storage and computation.

Our project fits this trend by using **Spark on EMR and S3** to manage large-scale preprocessing, and then applying **LSTM/GRU** for forecasting.

8 Data Preprocessing

Data preprocessing transforms raw S3-stored parquet files into model-ready datasets through cleaning, transformation, and feature engineering.

Data Cleaning and Formatting

Key cleaning steps include:

1. Timestamp conversion

All time columns are converted into Spark's `timestamp` type with a consistent timezone (e.g., UTC). This ensures that time-based functions like `hour`, `dayofweek`, and windowing operations behave correctly.

2. Duplicate and invalid record removal

Any duplicate records (same region, same timestamp) are dropped. Records with missing or negative load values are flagged and either imputed (if clearly a measurement error) or removed.

3. Handling missing weather data

Missing values in **temperature**, **humidity**, and **wind speed** are substantial. The project uses a combination of:

- Forward fill and backward fill for short gaps.
- Group-wise median imputation for longer gaps.
- Sensible defaults for non-critical variables if needed.

4. Consistency checks

Checks are performed to ensure:

- No timestamps are out of expected range.
- Region labels are standardized (e.g., no inconsistent spelling).
- Load units are consistent (MW).

column	missing_count
timestamp	0
load_mw	0
load_scaled	0
temperature	665235
humidity	668752
wind_speed	665220
is_holiday	891496
is_weekend	0
lag_1h	5
lag_24h	120
lag_7d	840
rolling_mean_24h	5
temp_x_weekend	665235
temp_x_holiday	891496
region	0
year	0
month	0

Feature Engineering

Feature engineering is performed using PySpark's window functions and transformations:

1. Temporal Features

Using `hour(timestamp)`, `dayofweek(timestamp)`, `month(timestamp)`, etc., we derive:

- Hour-of-day (0–23).
- Day-of-week (1–7).
- Month and year.
- Weekend indicator (based on day-of-week).
- Holiday indicator (using a calendar mapping for public holidays).

2. Lag Features

To capture autocorrelation:

- `lag_1h`: load at time $t-1$.
- `lag_24h`: load at time $t-24$.
- `lag_7d`: load at time $t-168$ (7 days * 24 hours).

These are implemented via Spark window specs partitioned by region and ordered by timestamp.

3. Rolling Averages

To smooth noise and provide context:

- 24-hour rolling mean.
- Optional 7-day rolling mean.

4. Weather Interactions

We create interaction features like:

- `temperature × weekend`.
- `temperature × holiday`.

These help capture the fact that temperature's impact is different on weekends or holidays compared to regular weekdays.

9 Tools Used

Programming Language

- **Python** is the primary language, chosen for its rich data science ecosystem and compatibility with Spark and TensorFlow.

Data Manipulation Libraries

- **PySpark** for distributed DataFrame operations, SQL-style queries, and window functions.
- **Pandas** for in-memory operations when converting Spark outputs to NumPy arrays for TensorFlow.

Visualization Libraries

- **Matplotlib** for static plots such as line graphs and histograms.
- **Seaborn** for aesthetically pleasing statistical visualizations (e.g., correlation heatmaps).

Imputation & Model Libraries

- **PySpark** and **Pandas** for custom imputation logic.
- **Scikit-learn** for scaling (e.g., `MinMaxScaler`) and possibly baseline models.
- **TensorFlow / Keras** for defining and training LSTM and GRU models.

Data Cleaning Tools

- Spark's DataFrame API and SQL functions are the main mechanisms for cleaning.
- Pandas is used for smaller, derived datasets (e.g., aggregated statistics or evaluation samples).

10 Exploratory Data Analysis (EDA)


EDA reveals patterns that guide feature engineering and model design.

Temporal Analysis

Load by Month (Seasonal Analysis)

A monthly aggregation of load (e.g., average load per month over several years) shows how consumption varies seasonally:

- **Summer months** (June–August) often display high loads due to air conditioning.
- **Winter months** (December–February) may also show elevated loads due to heating.
- **Spring and fall** can have moderate loads.




This confirms strong **seasonality**, which models need to capture.

Load by Day of the Week

Aggregating by day-of-week reveals:

- **Higher loads on weekdays**, when businesses and industries operate.
- **Lower loads on weekends**, reflecting reduced commercial and industrial usage.




This validates using **day-of-week** and **weekend indicators** as important features.

Load by Hour of the Day

Plotting average load for each hour of the day shows a typical pattern:

- Low during late-night hours.
- Increases in the morning as people wake up and businesses open.
- Peaks in late afternoon or early evening.
- Gradual decline at night.



This **daily cycle** explains the usefulness of `lag_1h` and `lag_24h`.

EDA computations use Spark group-by and aggregation operations, converted to Pandas for plotting. For example:


- `df.groupby("hour").agg(F.avg("load_mw"))`
- `df.groupby("dow").agg(F.avg("load_mw"))`

```
+-----+-----+
|timestamp|region|load_mw|
+-----+-----+
|2015-05-01 00:00:00|DOM|8084.0|
|2015-05-01 01:00:00|DOM|7548.0|
|2015-05-01 02:00:00|DOM|7293.0|
|2015-05-01 03:00:00|DOM|7182.0|
|2015-05-01 04:00:00|DOM|7210.0|
+-----+-----+
only showing top 5 rows
```

```
+-----+-----+-----+-----+-----+-----+
|region|avg_load|std_load|min_load|max_load|n_rows|
+-----+-----+-----+-----+-----+-----+
|AEP|10542.190280426248|7538.491735393321|0.0|9998.0|178300|
|COMED|4259.146690970275|5699.1173791303845|0.0|9999.0|178300|
|DAYTON|1386.0930846887268|1004.325613730939|0.0|982.0|178300|
|DOM|7135.036567582726|5568.806946606426|0.0|9999.0|178300|
|PJM|5491.9706331045|11816.464110727695|0.0|54030.0|178296|
+-----+-----+-----+-----+-----+-----+
```

Region vs Time Relationships


We can also examine how seasonality differs by region (e.g., some regions may be more summer-peaking, others more winter-peaking).



Correlation Analysis

A correlation heatmap (Pandas + Seaborn) displays correlations between:

- Load and lag/rolling features.
- Load and weather variables.
- Strong positive correlation between load and its own lagged values.
- Moderate correlation between load and temperature.
- Weaker correlation with humidity and wind speed.



Weather Correlation

More detailed plots (e.g., scatter plots of load vs temperature) show:

- Higher loads at extreme cold and extreme hot temperatures.
- Nonlinear shape (U-shaped or V-shaped), which deep learning can model.

Time and Location Relationship

Correlation between region/time and load can be analyzed by:

- Comparing patterns across regions.
- Checking if some regions are more weather-sensitive than others.

11 Predictive Modeling

This section mirrors the reference structure (KNN, Random Forest) but replaces them with **LSTM [1]** and **GRU** models.

LSTM Objectives

The LSTM model is designed to:

- Predict **next-hour load** based on a **24-hour input sequence** of load, weather, and temporal features.
- Capture both short-term correlations (last few hours) and daily patterns (yesterday's profile).

Preprocessing

Steps for LSTM preprocessing:

1. Select relevant features: load_mw, lag features, weather, temporal indicators.
2. Scale features using MinMaxScaler (fit on training data only).
3. Transform the time-ordered data into overlapping sequences of length 24.
4. Split sequences into training, validation, and test sets (e.g., 70/15/15).

Implementation

The LSTM model typically consists of:

- Input layer: shape (sequence_length=24, num_features).
- LSTM layer with 32 units and dropout.
- Dense hidden layer (optional) with ReLU activation.
- Output layer with a single neuron predicting the next-hour load (or scaled load).

The model is compiled with:

- Loss: Mean Squared Error (MSE).
- Optimizer: Adam.
- Metrics: MAE (for monitoring).

Callbacks used:

- EarlyStopping (patience ~3–5 epochs).
- ReduceLROnPlateau to lower the learning rate when validation loss plateaus.

Results

The LSTM model achieves approximately:

- MAE $\approx$ 550.65 MW
- RMSE $\approx$ 691.25 MW

```
csp554 — hadoop@ip-172-31-17-203:~/electricity_forecast — ssh -i emr-key-pair.pem hadoop@ec2-18-118-110-20.us...

or placeholder tensor 'gradients/split_grad/concat/split/split_dim' with dtype int32
[[[{{node gradients/split_grad/concat/split/split_dim}}]]]
2025-12-09 14:46:39.661290: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU:0] (DEBUG INFO) Executor sta
rt aborting (this does not indicate an error and you can ignore this message): INVALID_ARGUMENT: You must feed a value f
or placeholder tensor 'gradients/split_1_grad/concat/split_1/split_dim' with dtype int32
[[[{{node gradients/split_1_grad/concat/split_1/split_dim}}]]]
274/274 [=====] - 7s 17ms/step - loss: 0.2656 - mae: 0.3852 - val_loss: 0.0885 - val_mae: 0.236
0 - lr: 0.0010
Epoch 2/20
274/274 [=====] - 4s 16ms/step - loss: 0.1557 - mae: 0.2993 - val_loss: 0.0723 - val_mae: 0.214
4 - lr: 0.0010
Epoch 3/20
274/274 [=====] - 5s 19ms/step - loss: 0.1367 - mae: 0.2807 - val_loss: 0.0652 - val_mae: 0.202
5 - lr: 0.0010
Epoch 4/20
274/274 [=====] - 5s 20ms/step - loss: 0.1240 - mae: 0.2665 - val_loss: 0.0655 - val_mae: 0.203
4 - lr: 0.0010
Epoch 5/20
274/274 [=====] - 4s 16ms/step - loss: 0.1120 - mae: 0.2517 - val_loss: 0.0615 - val_mae: 0.195
9 - lr: 0.0010
Epoch 6/20
274/274 [=====] - 4s 16ms/step - loss: 0.1028 - mae: 0.2408 - val_loss: 0.0750 - val_mae: 0.218
8 - lr: 0.0010
Epoch 7/20
274/274 [=====] - 4s 16ms/step - loss: 0.0974 - mae: 0.2327 - val_loss: 0.0729 - val_mae: 0.213
7 - lr: 0.0010
Epoch 8/20
274/274 [=====] - 4s 16ms/step - loss: 0.0918 - mae: 0.2258 - val_loss: 0.0844 - val_mae: 0.232
1 - lr: 5.0000e-04
Evaluating LSTM fast...
2025-12-09 14:47:12.962987: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU:0] (DEBUG INFO) Executor sta
rt aborting (this does not indicate an error and you can ignore this message): INVALID_ARGUMENT: You must feed a value f
or placeholder tensor 'gradients/split_2_grad/concat/split_2/split_dim' with dtype int32
[[[{{node gradients/split_2_grad/concat/split_2/split_dim}}]]]
```

These metrics indicate the average absolute and squared deviations between predicted and actual load. Given the scale of the system, an RMSE of around 700 MW can be considered acceptable and competitive with similar studies in the literature.

Model: "LSTM_fast"

Layer (type)	Output Shape	Param #
input_layer_2 (InputLayer)	(None, 24, 13)	0
lstm_1 (LSTM)	(None, 32)	5,888
dense_4 (Dense)	(None, 16)	528
dropout_2 (Dropout)	(None, 16)	0
dense_5 (Dense)	(None, 1)	17

Total params: 6,433 (25.13 KB)
Trainable params: 6,433 (25.13 KB)
Non-trainable params: 0 (0.00 B)

Model: "GRU_fast"

Layer (type)	Output Shape	Param #
input_layer_3 (InputLayer)	(None, 24, 13)	0
gru_1 (GRU)	(None, 32)	4,512
dense_6 (Dense)	(None, 16)	528
dropout_3 (Dropout)	(None, 16)	0
dense_7 (Dense)	(None, 1)	17

Total params: 5,057 (19.75 KB)
Trainable params: 5,057 (19.75 KB)
Non-trainable params: 0 (0.00 B)

GRU

Objective

The GRU model is used as a **lighter alternative** to LSTM to explore:

- Whether a simpler recurrent architecture can achieve comparable performance.
- Slightly faster training times due to fewer parameters.

Preprocessing

Preprocessing is identical to LSTM:

- Same features.
- Same scaling.
- Same sequence construction and dataset splits.

Implementation

The GRU model:

- Replaces the LSTM layer with a GRU layer (e.g., 32 units).
- Keeps similar dropout and dense layers.
- Uses the same loss, optimizer, and callbacks.

Results

The GRU model achieves approximately:

- $MAE \approx 559.17$ MW
- $RMSE \approx 701.94$ MW


Model	MAE_MW	RMSE_MW
LSTM_fast	427.768310546875	573.6746137838069
GRU_fast	509.6291198730469	652.1008117998321

```
csp554 — hadoop@ip-172-31-17-203:~/electricity_forecast — ssh -i emr-key-pair.pem hadoop@ec2-18-118-110-20.us...
2025-12-09 14:45:31.537671: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU:0] (DEBUG INFO) Executor sta
rt aborting (this does not indicate an error and you can ignore this message): INVALID_ARGUMENT: You must feed a value f
or placeholder tensor 'gradients/split_1_grad/concat/split_1/split_dim' with dtype int32
[[[node gradients/split_1_grad/concat/split_1/split_dim]]]
274/274 [=====] - 7s 20ms/step - loss: 0.2257 - mae: 0.3593 - val_loss: 0.0740 - val_mae: 0.213
1 - lr: 0.0010
Epoch 2/20
274/274 [=====] - 5s 18ms/step - loss: 0.1385 - mae: 0.2814 - val_loss: 0.0604 - val_mae: 0.192
7 - lr: 0.0010
Epoch 3/20
274/274 [=====] - 5s 18ms/step - loss: 0.1218 - mae: 0.2627 - val_loss: 0.0563 - val_mae: 0.185
6 - lr: 0.0010
Epoch 4/20
274/274 [=====] - 6s 21ms/step - loss: 0.1086 - mae: 0.2461 - val_loss: 0.0509 - val_mae: 0.174
7 - lr: 0.0010
Epoch 5/20
274/274 [=====] - 5s 19ms/step - loss: 0.1007 - mae: 0.2365 - val_loss: 0.0541 - val_mae: 0.181
7 - lr: 0.0010
Epoch 6/20
274/274 [=====] - 5s 18ms/step - loss: 0.0950 - mae: 0.2279 - val_loss: 0.0472 - val_mae: 0.167
0 - lr: 0.0010
Epoch 7/20
274/274 [=====] - 5s 18ms/step - loss: 0.0892 - mae: 0.2206 - val_loss: 0.0451 - val_mae: 0.162
7 - lr: 0.0010
Epoch 8/20
274/274 [=====] - 5s 19ms/step - loss: 0.0854 - mae: 0.2147 - val_loss: 0.0491 - val_mae: 0.168
5 - lr: 0.0010
Epoch 9/20
274/274 [=====] - 5s 18ms/step - loss: 0.0813 - mae: 0.2085 - val_loss: 0.0444 - val_mae: 0.161
6 - lr: 0.0010
Epoch 10/20
206/274 [=====>.....] - ETA: 1s - loss: 0.0771 - mae: 0.203625/12/09 14:46:15 INFO YarnClientSchedule
rBackend: Requesting to kill executor(s) 2, 3
25/12/09 14:46:15 INFO YarnClientSchedulerBackend: Actual list of executor(s) to be killed is 2, 3
```

This is slightly worse than LSTM but still strong. The difference suggests that LSTM’s additional gating structure may help for this dataset.

LSTM vs GRU

- Comparison bar chart: LSTM vs GRU MAE / RMSE.
- Overlaid prediction vs actual curves.



12 Big Data Aspect

Volume

The dataset is truly “big” in the context of this course:

- Millions of rows (hourly records).
- Multiple regions and years.
- Many engineered features (lags, rolling windows, weather interactions).

Processing all these features in one pass on a single machine would be slow and memory-intensive; Spark and EMR solve this.

Variety

The data exhibits **variety** in several dimensions:

- Time-series load values.
- Weather features.
- Calendar-derived features.
- Region labels.

All are combined into a single feature set fed to models.

Tools and Techniques

To handle volume and variety:

- **S3** serves as the centralized storage layer.
- **EMR** provides a managed Spark cluster.
- **PySpark** handles feature engineering and EDA at scale.
- **EC2 + TensorFlow** trains deep learning models using data exported from Spark.

This demonstrates a **full big-data analytics pipeline**, not just a model in isolation.

13 Results

This section pulls together and interprets the main empirical findings.

LSTM Results

The LSTM model:

- Achieves lower MAE and RMSE than GRU.
- Tracks the load curve well in most periods.
- Tends to show larger errors during sudden spikes (e.g., extreme weather days).

An error analysis over time can reveal which days or hours contribute most to error (e.g., holiday peaks).

GRU Results

The GRU model:

- Slightly underperforms LSTM but still produces useful forecasts.
- May converge faster during training due to fewer parameters.
- Can be useful when computational resources are limited.

Comparing both helps justify the final choice of model for deployment (LSTM in this case).

14 Challenges Faced

Major challenges included:

1. **Weather Missingness**
Large fractions of weather data were missing, requiring careful imputation and potentially affecting forecast accuracy.
2. **Data Scale and Complexity**
The combination of multiple regions, years, and engineered features required efficient, distributed processing.
3. **Cloud Resource Constraints**
SageMaker Studio notebook instances were not available for the desired instance types/capacity, so the team adapted by using EMR and EC2 for training.
4. **CPU-Only Deep Learning**
Training LSTM/GRU models on CPU-only nodes is slower, requiring smaller model architectures and careful use of callbacks to manage training time.

15 Conclusion and Future Work

Key Findings

- An end-to-end **big data + deep learning pipeline** can be successfully implemented using S3, EMR, PySpark, EC2, and TensorFlow.
- **LSTM models** provide strong performance for short-term load forecasting, outperforming GRU in this project's configuration.
- Temporal and seasonal patterns, combined with weather information and lag features, are crucial for accurate load forecasting.

Overall Impact

Although developed in an academic setting, the architecture and methods used are representative of **real-world industrial forecasting systems**. The project demonstrates that graduate students can design and implement a sophisticated, cloud-based forecasting pipeline that integrates:

- Big data tools.
- Advanced modeling.
- Practical constraints and adaptation.

Model Improvement

Future model improvements could include:

- Incorporating external, high-resolution weather data (e.g., from APIs) to reduce missingness and improve feature quality.
- Adding **additional exogenous features**, such as price signals, renewable generation forecasts, or demand response schedules.
- Using **multi-output models** to predict multiple hours into the future.

Advanced Techniques

Advanced directions:

- Applying **Transformer-based architectures** for long-sequence modeling.
- Using **probabilistic forecasting** (e.g., quantile regression, Bayesian neural networks) to provide prediction intervals.
- Deploying the model with real-time inference using managed services like SageMaker endpoints or custom EC2 APIs.

16 Acknowledgment

The team would like to thank:

- The **CSP 554 course instructors** for guidance and for designing a project that bridges big data technologies and machine learning.
- The providers of the **load and weather datasets**, whose publicly available data made this study possible.

References

- [1] S. Hochreiter and J. Schmidhuber, “Long Short-Term Memory,” *Neural Computation*, vol. 9, no. 8, pp. 1735–1780, 1997.
- [2] G. Box, G. Jenkins, and G. Reinsel, *Time Series Analysis: Forecasting and Control*, 3rd ed. Englewood Cliffs, NJ: Prentice-Hall, 1994.
- [3] K. Cho et al., “Learning Phrase Representations using RNN Encoder–Decoder for Statistical Machine Translation,” in *Proc. EMNLP*, 2014, pp. 1724–1734.
- [4] R. J. Hyndman and G. Athanasopoulos, *Forecasting: Principles and Practice*, 2nd ed. OTexts, 2018.
- [5] D. Marino, K. Amarasinghe, and M. Manic, “Building Energy Load Forecasting Using Deep Neural Networks,” in *Proc. IEEE IECON*, 2016, pp. 7046–7051.
- [6] W. Kong et al., “Short-Term Residential Load Forecasting Based on LSTM Recurrent Neural Network,” *IEEE Trans. Smart Grid*, vol. 10, no. 1, pp. 841–851, 2019.
- [7] PJM Interconnection, “Hourly Load Data,” 2023.
- [8] NOAA National Centers for Environmental Information, “Global Historical Climatology Network (GHCN),” 2023.
- [9] M. Zaharia et al., “Apache Spark: Cluster Computing with Working Sets,” in *Proc. USENIX HotCloud*, 2010.
- [10] Apache Spark Project, *MLlib: Scalable Machine Learning Library*, 2023.
- [11] M. Abadi et al., “TensorFlow: Large-Scale Machine Learning on Heterogeneous Systems,” in *Proc. OSDI*, 2016.
- [12] Amazon Web Services, *Amazon SageMaker Developer Guide*, 2023.